



# Introduction to Machine Learning & Generative AI

A Technical Deep Dive into Predictions, Neural Networks, and Modern AI

AI/ML FUNDAMENTALS

MENTOR DISCUSSION

# Session Agenda

01

## Machine Learning Foundations

Core concepts and real-world importance

03

## Neural Networks & Perceptrons

Architecture, training, and key differences

02

## Prediction & Regression

How models make forecasts and learn patterns

04

## Generative AI

Modern applications and transformative impact



# What is Machine Learning?

## Definition

Machine Learning enables computers to learn patterns from data without explicit programming, improving performance through experience.

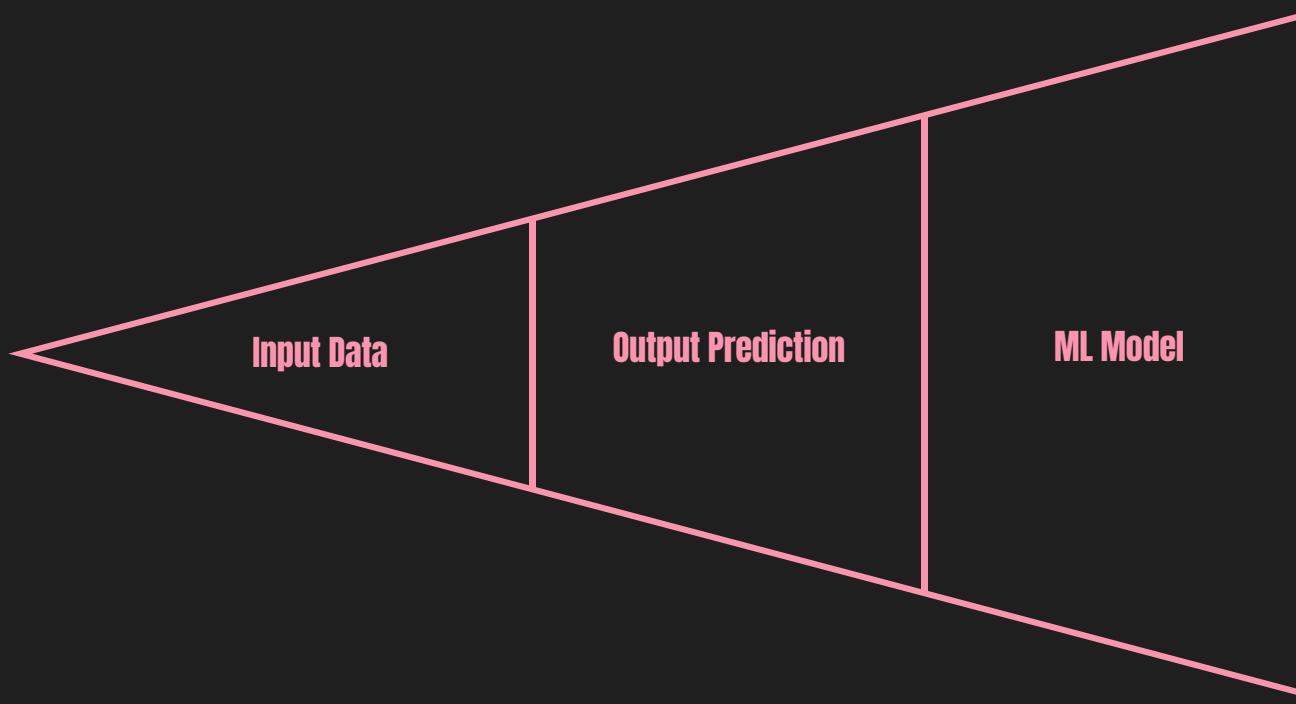
## Why It Matters

ML powers everything from medical diagnosis to autonomous vehicles, transforming industries by uncovering insights humans might miss.

## Real-World Example

Netflix's recommendation engine analyzes viewing habits to suggest content, keeping users engaged through predictive algorithms.

# Prediction in Machine Learning



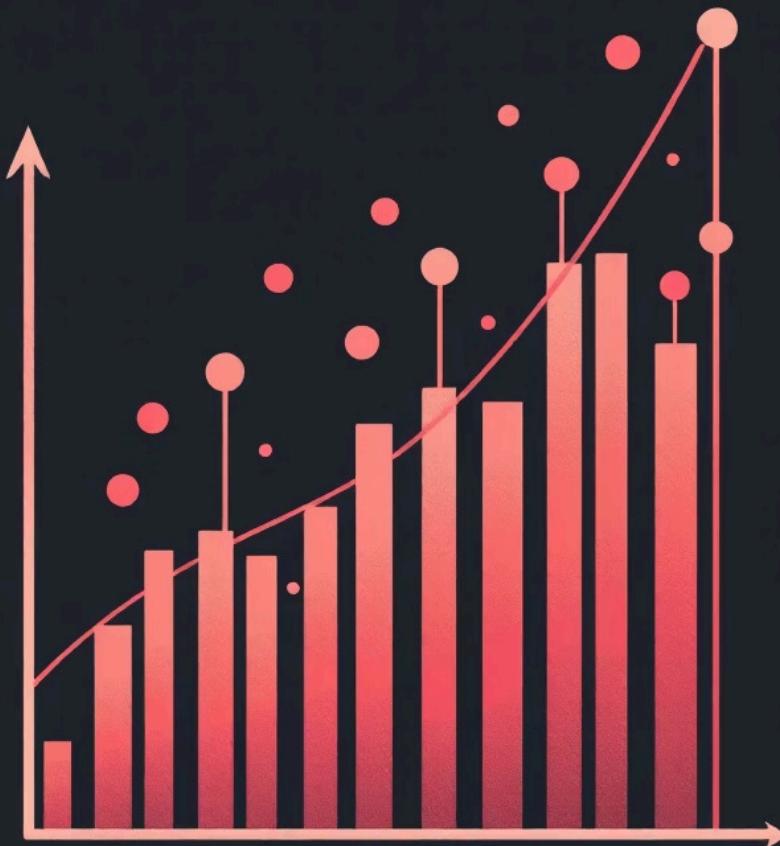
## How Prediction Works

Prediction is the process where a trained model takes new input data and generates an output based on learned patterns.

- Input: Features like age, income, or past behavior
- Model: Trained algorithm that processes inputs
- Output: Predicted value or classification

**Example:** Predicting house prices based on location, size, and amenities using historical sales data.

# Understanding Regression



## Linear Regression

Models the relationship between a single input variable and output using a straight line. Perfect for simple cause-effect relationships.

## Multiple Regression

Extends linear regression to handle multiple input variables simultaneously, capturing complex real-world scenarios.

## Common Use Cases

Sales forecasting, stock price prediction, risk assessment, customer lifetime value estimation, and demand planning.

# Regression in Action

## The Formula

$$y = mx + c$$

Where:

- **y** = predicted output
- **m** = slope (weight)
- **x** = input feature
- **c** = y-intercept (bias)

## Python Implementation

```
from sklearn.linear_model import LinearRegression

# Training data
X = [[1], [2], [3], [4], [5]]
y = [2, 4, 6, 8, 10]

# Create and train model
model = LinearRegression()
model.fit(X, y)

# Make prediction
prediction = model.predict([[6]])
print(prediction) # Output: [12]
```

# Neural Networks Fundamentals



**Artificial Neuron**  
Mimics biological neurons by receiving inputs, applying weights, and producing an output through an activation function.



**Layer Architecture**  
Networks consist of input layers (receive data), hidden layers (extract features), and output layers (make predictions).

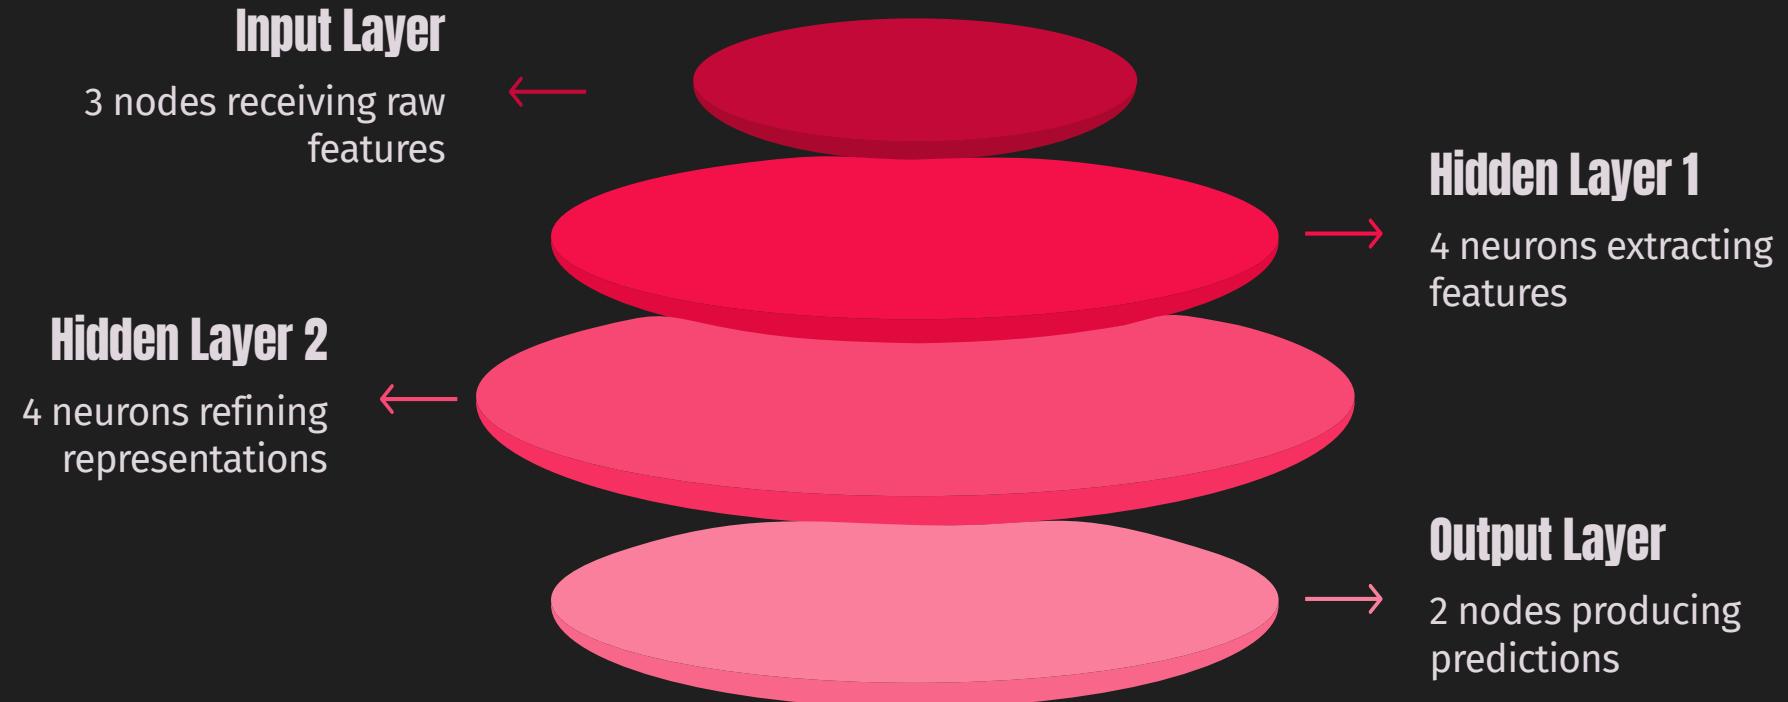


**Activation Functions**  
Non-linear functions like ReLU, sigmoid, and tanh introduce complexity, enabling networks to learn intricate patterns.



**Training Process**  
Networks learn through backpropagation, adjusting weights iteratively to minimize prediction error using gradient descent.

# Neural Network Architecture



# The Perceptron Model



## Simplest Neural Unit

The perceptron is a single-layer neural network that performs binary classification by learning a linear decision boundary.

## How It Works

1. Receives multiple weighted inputs
2. Sums inputs:  $\sum(w_i \cdot x_i) + b$
3. Applies activation function (step function)
4. Outputs 0 or 1 based on threshold

## Core Equation

$$y = \begin{cases} 1 & \text{if } w \cdot x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

# Perceptron vs Neural Network

Aspect	Perceptron	Neural Network
Layers	Single layer	Multiple layers (deep)
Complexity	Linear problems only	Non-linear patterns
Activation	Step function	ReLU, sigmoid, tanh
Learning	Simple weight updates	Backpropagation
Use Case	Basic classification	Image, text, speech recognition

- ❑ **Key Limitation:** Perceptrons cannot solve problems that aren't linearly separable, like XOR. Neural networks overcome this with hidden layers.